

RUTGERS UNIVERSITY

CAPSTONE SENIOR DESIGN PROJECT

14:332:492

Scarletshield

A Lightweight, Linux-based Network Security Solution Suite

Team:

Jeff ADLER

Eric CUIFFO

Parth DESAI

Jeff RABINOWITZ

Val A. RED

Advisor:

Dr. Manish PARASHAR

April 30, 2014

Contents

1	Abstract	2
1.1	Introduction	2
2	Hardware and Software Architecture	5
2.1	Evaluation of System Hardware	5
2.2	System Software Architecture	6
3	Approach and Methodology	7
3.1	Approach	7
3.2	Securing SSH	7
3.3	Installation	8
4	Administration Suite & Features	10
4.1	Web Front-end Integration	10
4.2	Analytics	11
4.3	Response Options	13
5	Conclusion	15
5.1	Future Work	15

Chapter 1

Abstract

“Scarletshield” is a customized Ubuntu 12.04 LTS image featuring a uniquely layered cyber security deployment of several open source services; most notably the **”Snort” Intrusion Detection System (IDS)** utilizing packet inspection and the **”Fail2ban” Intrusion Prevention System (IPS)** utilizing log inspection. It is optimized for synergetic network defense and complemented by a dynamic, modern web frontend providing the utmost flexibility for network administrators to monitor and quickly react to any and all threats against a network. The design of was conceived as a proof-of-concept to expand upon the more orthodox but aging Defense in Depth (DiD) strategy that layers *overlapping* technologies rather than presenting varied layers of *complement* defense services working in *synergy*, which is the approach use to expand upon the defense-in- depth strategy. In addition, Scarletshield enables the potential for a user to interface with interconnected subnets and gateways over a private network to enable scalability and mitigate large-scale attacks. Overall, Scarletshield is intended to be a flexible, open-source system for preventing and thwarting evolving Distributed Denial-of-service (DDoS) Attacks and Advanced Persistent Threat (APT) by implementing a synergy of various open source and custom designed services designed to be as flexible and as intuitive as possible for the network administrator to interface with, minimizing administrative overhead and maximizing mobility for the network administrator to track and react to threats.

1.1 Introduction

Attackers have the edge in the dynamically evolving field of cyber security. With a wider breadth of strategies, near-infinite resources, and a virtually untraceable mask of anonymity; hackers have long boasted the advantage, successfully rendering useless the generally accepted DiD layered defense mechanisms most organizations deployed for network defense. Essentially, is analogous to an onion: although it is deeply layered, it can be peeled. Utilizing an even wider breadth of tools and resources, hackers have successfully and continuously

developed new and advanced approaches, exploits, Denial-of-service (DoS) methods, etc. to employ a persistent, peeling threat designed to break down or even bypass the DiD approach employed by network administrators. In essence, this summarizes the APT, which is not a single exploit/attack or even a collection of exploits/attacks: APT is a *campaign* involving possibly several of different collections of exploits and attack methodologies and even multiple actors across different machines that may exist in different countries.

Defense in Depth (DiD) [1], as described by the National Security Agency (NSA), is “[a] practical strategy for achieving Information Assurance in today’s highly networked environments. It is a “best practices” strategy in that it relies on the intelligent application of techniques and technologies that exist today.”

While DiD appears sound in its introduction, a closer and modern interpretation of later details in the NSA strategy employing defense-in-depth shows some of its flaws that come with age, for example: “To effectively resist attacks [...] an organization needs to *characterize its adversaries*, their potential motivations, and their *classes of attack*.”

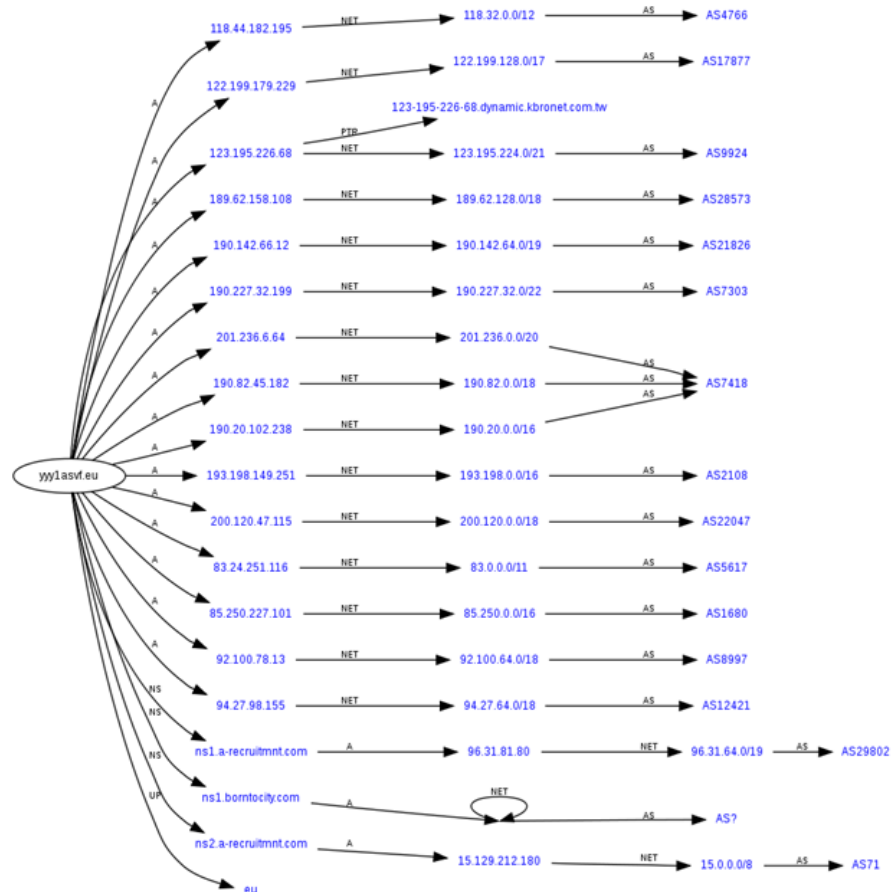


Figure 1.1: A visualization of the domain fast-flux. Note the large number of IP addresses and domains mapping to a single fully qualified domain name used for the C2 server.

The above notion presented by the NSA is idealistic at best in today’s climate of virtually

untraceable, anonymous hackers employing robust Domain Name Service (DNS) techniques such as **fast flux** [2], which essentially accomplishes hiding a single Command-and-control (C2) server by utilizing numerous IP addresses and a single, rotatable (in an even more difficult-to-trace **double-flux** methodology) fully qualified domain name (something as simple as valred.com or hackr.ru, for example) to very quickly swap between the available IP addresses and DNS records to avoid detection and backtracing. Thanks to these particular methods, problems such as phishing will be omnipresent in the foreseeable future, exasperated by other countries boasting lax restrictions with regulating the domain name provisioning. As such, the NSA’s notion of having organizations “characterize its adversaries,” is among the highly impractical and dated points that weaken the argument for the defense-in-depth strategy. The only likely moment adversaries are to be characterized is when it is too late and a network has already been compromised, such as when the Syrian Electronic Army (SEA) successfully executed a man-in-the-middle attack rerouting traffic to the U.S. Marines [3] in 2013. As such, we keep Scarletshield lightweight and practical by not addressing this issue; instead, we specifically handle threat prevention and detection.

Strengths of the defense-in-depth methodology are applied in the implementation of Scarletshield . Particularly, one directed focus is on protecting the most exposed web-facing systems that are typically the target of automated attacks. This is specifically handled by the robustness of the Snort IDS, which essentially sniffs all types of Internet Protocol (IP) packets (TCP, UDP, ICMP, etc.), and checks against patterns, or *rules* and *signatures* indicative of an attack. This is reinforced even further by the log-inspection capabilities of fail2ban, which can essentially pattern-match very common exploits and DoS traces against typical protocols such as HTTP, SSH, and FTP and react by blocking IP addresses logged via “iptables”, a powerful firewall built into most modern Linux distributions (distros) .

Even though iptables comes with modern Linux distros by default, many end users beyond actual system and network administrators fail to utilize iptables to its full potential. Scarletshield can partly fill the role for a network administrator when an end user does not have the personnel or resources to fully utilize iptables themselves. It can essentially be connected to any switch or router and be configured to act as a network gateway such that all traffic to any end user on a network will have to make it past Scarletshield first.

One issue that comes with defense-in-depth is overhead. How will a user interface with Snort and fail2ban? Scarletshield bridges the gap by presenting all IDS and IPS information and actions in a modern, sleek front-end accessible via private network (so only computers in the Scarletshield network may access it) in a way that minimizes the overhead of a network or system administrator having to manage everything via command line.

In summary, Scarletshield employs a depth of security systems including Snort and fail2ban while also boasting a breadth of options such as iptables handling and dropping of malicious sources’ packets while interfacing with other gateways and subnets to maximize overall security in a network.

Chapter 2

Hardware and Software Architecture

2.1 Evaluation of System Hardware

Scarletshield in its current manifestation runs in the Microway server described below. Note how the hardware is actually dated by today's standards; this exemplifies the scalability of Scarletshield – it does not actually require a high-end server to operate efficiently, but can actually operate on smaller devices as well. A more detailed study of scalability factors occurs in Chapter 4.

The largest performance factor when deploying the server for monitoring gateways is the Network Card. In the small, local, web-facing network we employed, a one-gigabit network card was sufficient enough to handle all requests and still have Scarletshield simultaneously replicate and check the requests against our IDS and IPS configuration. However, for maximum scalability, a network administrator should mind the network card being used, and if a Scarletshield-esque design is to be employed, also ensure that a second network card is present to facilitate DHCP and network connectivity for local machines within the network.

Chipset	ASUSTek Computer INC. K8N-DRE
CPU	2x Dual Core AMD Opteron(tm) 275 : 2.2GHz 2x1 MB L2 Cache
RAM	4x8GB
OS	Ubuntu 12.04 LTS
HDD	60GB HDD
Network Card	2x NetXtreme BCM5721 Gigabit Ethernet PCI Express

Table 2.1: The hardware specifications for the Microway Server provided by Rutgers Engineering Computing Services (ECS).

2.2 System Software Architecture

As shown below, Scarletshield is composed of two interfacing subsystems – a packet inspection and logging system, composed of Snort, Barnyard2, Pulled Pork, fail2ban, and iptables; and a web-based administrator suite powered by Drupal. These two systems are married by a MySQL server, which holds both Snort’s configurations, as well as the attack logs.

The Administration Suite utilizes Drupal, a Content Management System (CMS), not only for its maturity, but also for its modularity and extensibility; PHP scripts can be added to Drupal’s configuration and then served to web clients with minimal effort. This is actually how the Scarletshield Administration Suite operates – through a series of modular PHP scripts furnished to Drupal. This minimalistic approach has the advantage that adding a new page to the suite is as simple as writing another script and providing it to Drupal. In this way, Scarletshield can be customized by end users to meet individualized needs.

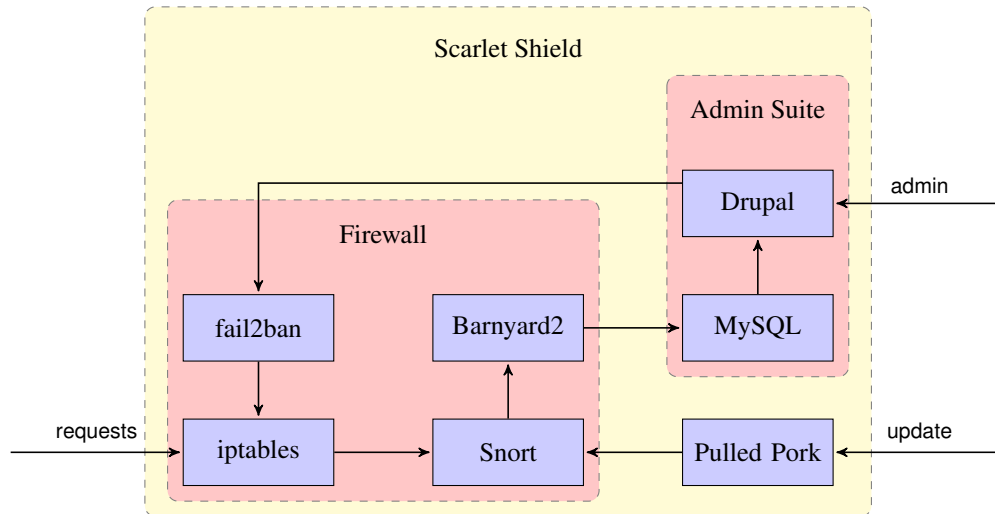


Figure 2.1: The architecture of Scarletshield. Incoming packets are routed through iptables, where they may be blocked per iptables’ rule-set; Snort receives the packets and inspects them for signatures of attacks; offending packets are mirrored via Barnyard2 into a MySQL database; users can access the Administrator Suite provided by Scarletshield and hosted using Drupal in order to analyze attacks against the system; fail2ban automatically bans abusive foreign hosts. Pulled Pork is used to automatically update Snort’s classification rulesets for offending packet patterns.

Chapter 3

Approach and Methodology

As we are designing Scarletshield from scratch with a Ubuntu 12.04 LTS server, there are several sections to our Approach and Methodology:

1. Approach
2. Secure Server SSH
3. Set up of Snort and fail2ban

3.1 Approach

Scarletshield takes all incoming packets and either accepts or drops them based on iptables, which is built in to modern Linux distros such as Ubuntu 12.04 LTS automatically. What makes Scarletshield unique is how we deploy an IDS and IPS through Snort and Fail2ban. Specifically, IDS mirrors all packets and analyzes them for patterns (via Snort rules, which are updated every month over the Internet) indicative of a potential threat, then logging threat signatures on a database accessible to the administrator and analyzed for new, more strict iptables rules banning offending IP addresses where necessary. We utilize fail2ban to detect obvious brute force, overflows, and exploits in the server logs (/var/log/auth.log, etc.) to also automatically filter and block IP addresses preventively, requiring no action from the administrator.

3.2 Securing SSH

Through our research as a group, we came to the conclusion that the safest way to protect our SSH server against brute forcing would be to change our SSH's port from 22 to create security

through obscurity. However, in order to gather the largest possible set of logged attack data, we needed to present SSH on its standard port. This led us to a different approach of setting up two-factor authentication. By utilizing Google’s “Google Authenticator API”, in order to successfully gain access to ScarletShield’s SSH server, one must first enter the correct password and the unique TOTP security token, described by RFC6238 and generated by the “Google Authenticator” app, for the specific user trying to login. This security measure makes it impossible for a malicious user to gain login access through brute force.




Figure 3.1: (Left) A prompt to enter a Google Authenticator verification code after typing in a password. (Right) Google Authenticator smart phone application generating a verification code.

3.3 Installation

Scarletshield covers the breadth of IPS and IDS by implementing both fail2ban and Snort and interfacing the former with iptables, a user space application for packet handling. This not only enables the mobility for a network administrator to easily be able to monitor incoming traffic and potential threats on a switch, but also provides an automated mechanism via fail2ban to block obvious, trivial threats such as brute force attacks and denial-of-service.

Fail2ban is a very simple installation in Ubuntu 12.04 LTS, since it is included in the advanced packaging tool. Using its internal mechanism known as “jails”, it essentially utilizes filters implemented based on obvious threats (brute force over SSH, flooding over apache, etc.) and, over a user-set period of time (default period is several minutes long), bans offending IP addresses for the duration after it gets caught by Fail2ban via iptables directives packaged with the program.

Snort is a larger beast with many different variations of configurations available (types of relational databases used, logging mechanism, etc.); at the very least, an administrator separately needs to acquire and install libdnet and the Data Acquisition API [4]. In addition,

an administrator must choose a database and correctly install and configure its respective libraries and server application. Due to how fully featured it has become over the years, we chose MySQL.

Once Snort and the database being used are installed, the user must then acquire and/or update the Snort rule-set, which is essentially how Snort is able to parse and detect threats within the headers and contents of packets entering the machine. We chose “Pulled Pork”, as it is the most reliable Snort rule management application, recommended by the creators of Snort themselves. Finally, the administrator must interface Snort with the MySQL database. As of Snort 2.9.3.0, the latest stable version available for Ubuntu 12.04 LTS, it is required to install the application “Barnyard2” for a Snort daemon to record information into the MySQL database created by a user.

Chapter 4

Administration Suite & Features

4.1 Web Front-end Integration

Scarletshield features a web frontend utilizing a simple but powerful Drupal Content Management System (CMS). In addition, several PHP scripts are utilized to help analyze and serve the Snort MySQL database.

Scarletshield has three powerful tools for a user that wishes to learn more about the offending attacks against their server. These three tools, the “Attack History”, “Analytics” and “Heatmap” pages, all come comprise Scarletshield Suite. Each tool features a different way for the user to view the logs stored by Scarletshield. The link for the Scarletshield Suite is <http://scarletshield.rutgers.edu/suite>.¹

When a user tries to access any page of the suite, a JavaScript function runs to see if the user has an active login cookie with the correct password (the password is “scarlet”). If the user does not have this cookie, they are redirected to the suite login page seen below.

The login page makes use of two JQuery packages. The first package, “Tubular”, allows a streamed video to be the page background. The other package, “Toastr”, makes use of JavaScript’s non-blocking architecture. Toastr is responsible for all asynchronous pop-up alerts. These can be seen at page load, and when the access key form’s “Submit” button is pressed. Upon entering the correct password, the login cookie is generated and the user is then brought to the “Attack History” page. When first accessing this page, a loading animation is shown, while the back-end queries data from the MySQL database containing all Snort logs. As a large dataset is loaded from the server, a screen similar to the one shown below is displayed.

¹Note that the Scarletshield Suite was written for the Chrome browser, and some small details will not work on other browsers.




Figure 4.1: The Scarletshield Suite login page.

4.2 Analytics

Immediately upon visiting the page, the user is given access to all the data that Snort has logged; the data are logically and visually organized and to maximize presentability. (For comparison, an early implementation of this page can be seen here: <http://scarletshield.rutgers.edu/demo/last100.php>.) Each log is grouped by date and contains information

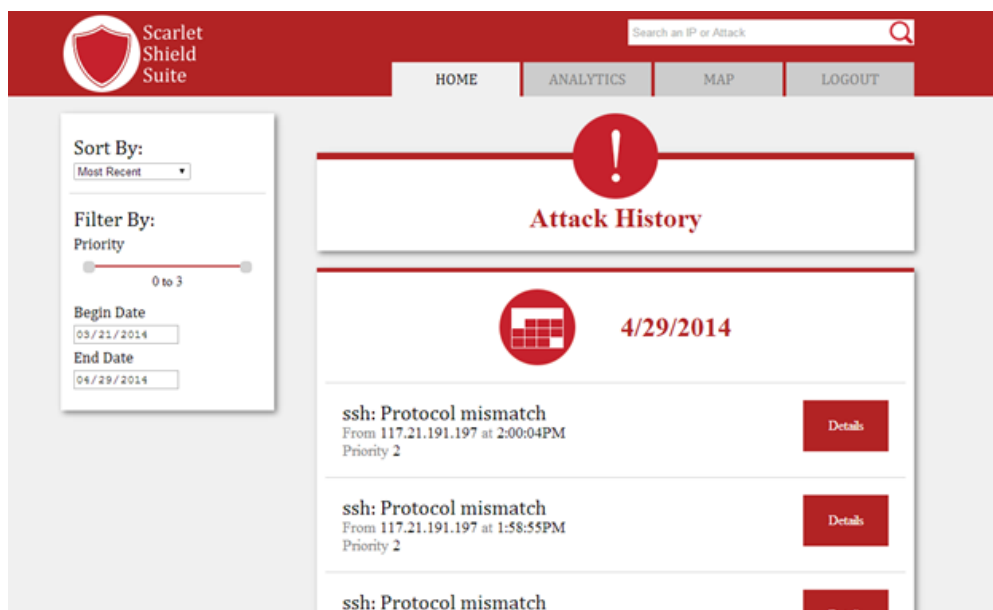


Figure 4.2: The Scarletshield Suite Attack History page.




Figure 4.3: The Scarletshield Suite Attack History page, with an expanded description of the details of an attack originating from China.

about the type of attack, the IP address it was originated from, the time it occurred and the priority of the attack (on a range from 0 to 3 with 0 being highest priority). If the user presses the “Details” button, extra information, such as the destination IP address (which may not be static if the server has more than one IP address), the ID number of the occurrence, and approximate location of the attack are expanded.

On the top left hand sides of the page, several tools allow the user advanced filtering capabilities. Logs can be sorted by date (oldest or newest) and priority (lowest first or highest first), and the search bar at the top can be used to only show information that either contains the IP address we search or show attacks that match the string that was searched. This is incredibly convenient because manually querying for signatures in MySQL (which involves several tedious SQL joins) would be a time-consuming task for an administrator, and by having it presented via an admin-accessible private frontend, not only could administrators react quicker, but even less-skilled users could discover and report an ongoing attack.

The Analytics page has three sections, as noted on the sidebar: the “Time Analysis”, “Common IP’s”, and “Common Attacks”. The Time Analysis section is simply a chart that shows a chart with the number of attacks that the server has undergone each day. This chart has a slider to focus on specific time intervals, and as one can see from the entire view, this is necessary because one single day may be so active that it dwarfs other days’ activity.

The last page in the Scarletshield Suite is the Heatmap page. Essentially, Scarlet Shield processes IP addresses from all logged attack patterns and places it on a Google Map heatmap, so that the user may be able to geographically backtrace signatures. In certain configurations, the heatmap would be useful to backtrace potential ongoing threats and

attacks, such as DoS attacks within certain windows. This would also help administrators make informed decisions for creating temporary iptables rules for banning certain offending subnet servers for the duration of an ongoing threat. An example of the heatmap is shown below.

4.3 Response Options

The final portion enabling Scarletshield to cover the full spectrum of network defense is enabling the network administrator to utilize the frontend for analysis and decision-making options. Essentially, the an administrator should be able to take the Snort records, make an analysis, and decide the action (whether to throttle or drop all packets altogether from a suspicious IP address) to take to ensure network security.

Due to time constraints, this feature was not integrated into the administrator suite, although the capability is already inherently present because of the inclusion of fail2ban. Future work could easily incorporate the addition of such an administrator response suite.

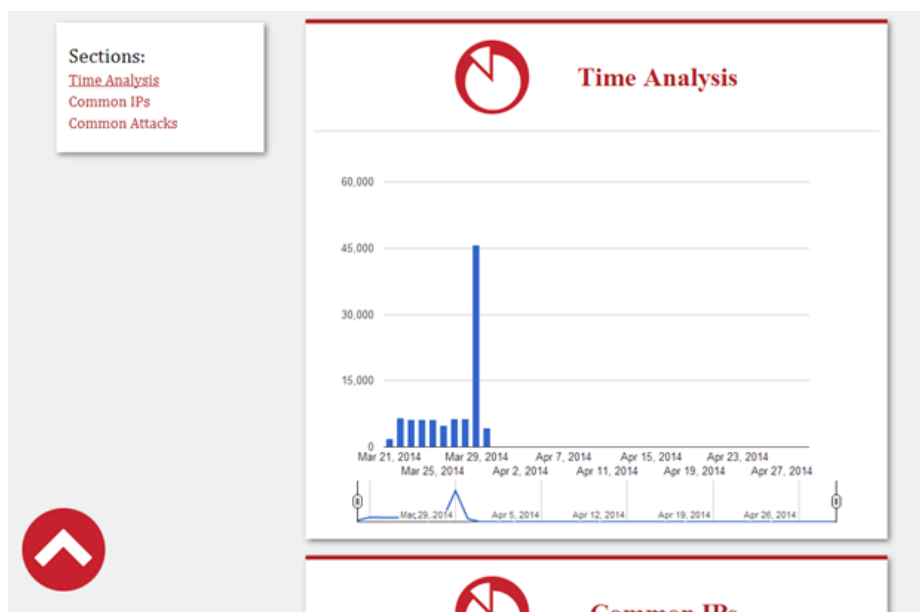


Figure 4.4: A chart of the rate of attacks over a certain interval. Users can use the slider at the bottom of the chart to analyze attacks over arbitrary intervals.


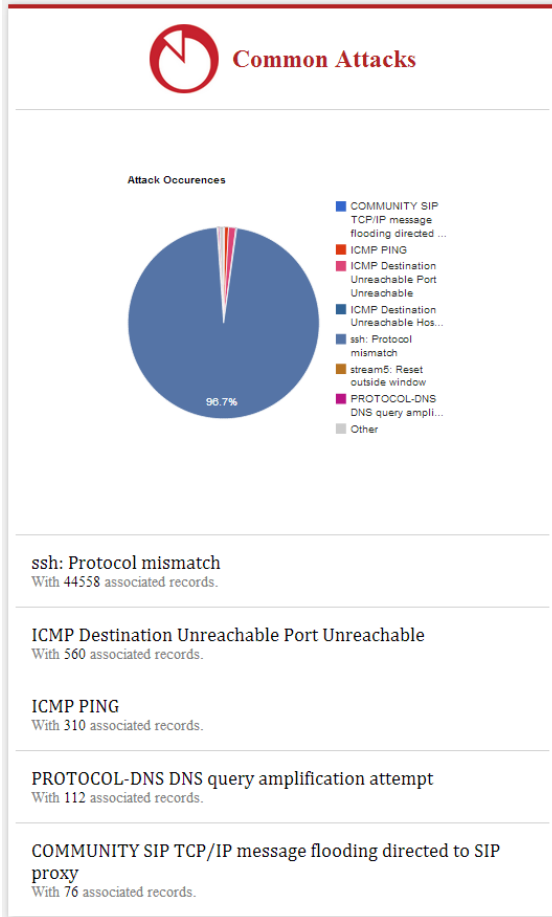


Figure 4.5: Pie charts featuring prominent attack varieties and IP addresses.

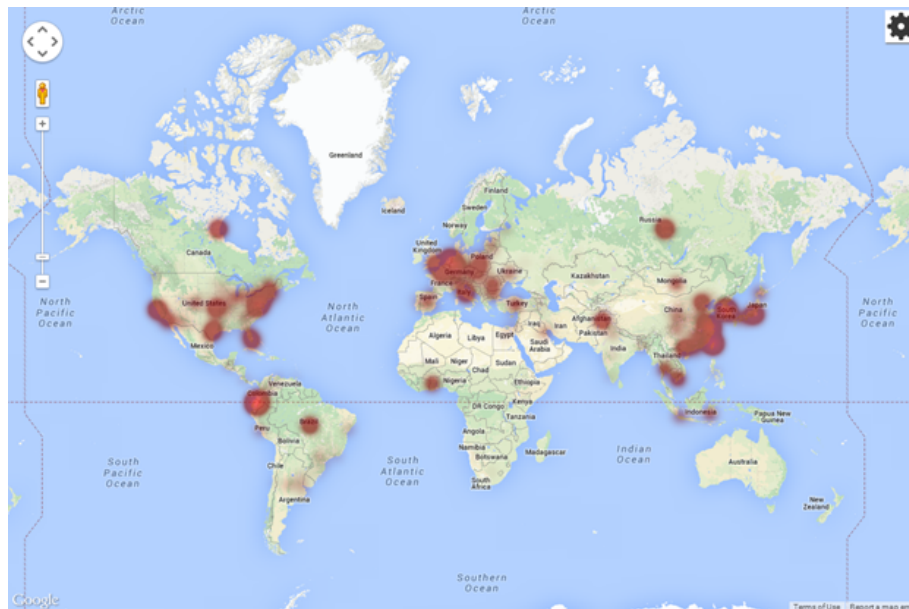


Figure 4.6: A geographic world map with red intensity markers indicating the presence of attack vectors from a geographic region, identified by IP address.

Chapter 5

Conclusion

During this semester, we achieved the following goals:

Our work in integrating the Snort IDS, fail2ban IPS, and iptables within the <http://scarletshield.rutgers.edu> web server has culminated in the final development and release of a robust frontend, the Scarletshield Suite, which is available via <http://scarletshield.rutgers.edu/suite> with the password “scarlet”. Our overall framework presents a great proof-of-concept for an open-source solution for network security that can be applicable to small-time enterprises and individuals who would not prefer to invest in a proprietary or commercial IDS or IPS.

To exemplify how effective our work in Scarletshield has been this semester, Rutgers Engineering Computing Services (ECS) has tasked system administrator Val A. Red with replicating the Scarletshield Suite for internal use to monitor wireless access point gateways.

5.1 Cost/Sustainability Analysis and Scalability

In spite of existing in an aged Microway rackmount server, the primary component necessary for utilizing Scarletshield on a network gateway is the Network Card. While a dedicated rackmount server is a costly investment, much of the hardware that Scarletshield actually runs off of is dated and can be purchased for very cheap prices off the shelf. One of its most important elements, a NetXtreme BCM5721 Gigabit Ethernet PCI Express card, actually goes for \$15-\$70 (depending on where you look, and on whether it is used or refurbished).

Scarletshield in its current manifestation operates off of a Microway rackmount server. An equivalent, but modern-day, rackmount server would range between \$2,000-\$5,000, which would be a large investment.

5.2 Future Work

Although Scarletshield is already operational on a server within Rutgers Engineering, additional work is necessary to integrate it into private subnets and to give it world-wide-web scale. More robustness and failover capabilities are needed in order to communicate certain threats that may be large-scale, and which would otherwise disrupt significant cybersecurity infrastructure. This also requires fine tuning Scarletshield's frontend capabilities and adding more analysis tools and automation techniques to streamline the threat-analysis and threat-prevention workflow presented to the network administrator. Further integration of the sub-components of Scarletshield will solidify its versatility, relevance, and overall strength. By successfully implementing a fully-featured frontend for monitoring and reacting to threats, we will maximize the potential of Scarletshield to be deployable to both private switches and enterprise-level networks.

Primarily, this entails better integrating current monitoring scripts with Drupal.

Moving forward we envision multiple ScarletShield deployments being used to protect a multitude of systems simultaneously. With this in mind we plan on creating a global set of black lists and logs that all deployments can access and update. Maintaining global lists would allow each individual deployment of ScarletShield to share collective knowledge. If one instance of ScarletShield is attacked, another deployment can preemptively block the attacker instead of waiting to be attacked itself.

Also, our scripts will allow administrators the option to take action on addresses logged with signatures, along with varying degrees of reaction steps they may be able to take (functionality to throttle or ban IPs altogether, for example), such that Scarletshield will give full control to the network administrator.

Additionally, we will be working on a way to replicate the Scarletshield build; this will most likely be accomplished by presenting a ISO image that can be installed in managed switch servers.

Bibliography

- [1] United States National Security Agency. (2012). Defense in Depth: A practical strategy for achieving Information Assurance in today's highly networked environments, [Online]. Available: http://www.nsa.gov/ia/_files/%20support/defenseindepth.pdf.
- [2] J. Riden. (2008). How fast-flux service networks work, [Online]. Available: <http://www.honeynet.org/node/132>.
- [3] A. Sternstein. (2013). Pentagon planned to secure web domains before syrians hijacked marine corps site, [Online]. Available: <http://www.nextgov.com/cybersecurity/2013/09/pentagon-planned-secure-web-domains-syrians-hijacked-marine-corps-site/69830/>.
- [4] D. Gullett. (Jul. 22, 2012). Snort 2.9.3 and Snort Report 1.3.3 on Ubuntu 12.04 LTS Installation Guide, [Online]. Available: <http://www.snort.org/assets/158/snortinstallguide293.pdf>.